

AWS Feedback Management System

Work Integrated Project

Karthikeyan Jeyabalasuntharam

Program: Cloud Computing Technologies (T465)

Group Name: CloudZen

Contents

Project Overview	3
Objectives	3
AWS Architecture Diagram	4
AWS Resource	6
Implementation Steps	7
Conclusion	37

Project Overview

The AWS Feedback Management System is a multi-tier cloud application designed to collect, process, analyze, and visualize feedback from users in real-time. It is deployed entirely on Amazon Web Services (AWS) using secure, scalable, and highly available architecture.

The system follows a three-tier architecture:

- Frontend (Presentation Layer) – Provides a user interface where users submit feedback through a form.
- Backend (Application Layer) – Processes incoming feedback, stores it in the database, and integrates with AWS Comprehend for sentiment analysis.
- Database (Data Layer) – Stores persistent feedback data, including sentiment analysis.

A key feature of this system is Auto Scaling, which automatically adjusts the number of application and web and app server instances based on incoming traffic. This ensures optimal performance during peak loads while reducing costs during low traffic periods.

In addition to collecting feedback, the system uses AWS Comprehend to detect sentiment (positive, negative, neutral). Monitoring, alerting, and logging are handled by Amazon CloudWatch, ensuring operational visibility and rapid troubleshooting.

Website Link - <https://karthikcode.com/>

Objectives

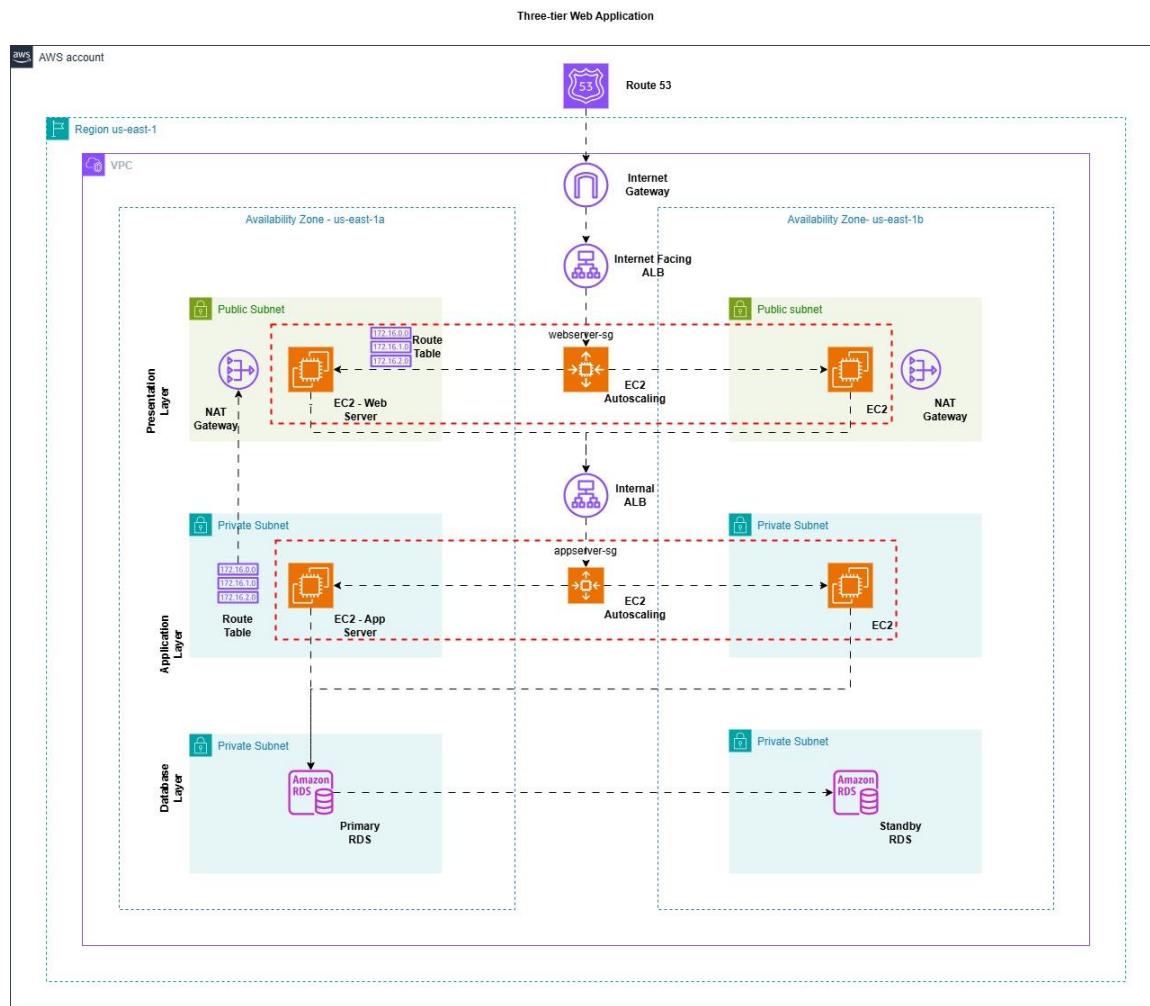
The objectives of the AWS Feedback Management System are:

- Collect User Feedback: Provide a simple and accessible web interface for feedback submission.
- Process and Store Feedback Securely: Ensure data integrity and confidentiality using Amazon RDS.
- Analyze Feedback Sentiment: Integrate AWS Comprehend to automatically classify feedback sentiment.
- Implement Auto Scaling: Dynamically scale EC2 instances for the web and app tiers to match workload demands.
- Maintain High Availability: Deploy resources across multiple Availability Zones with load balancing for fault tolerance.

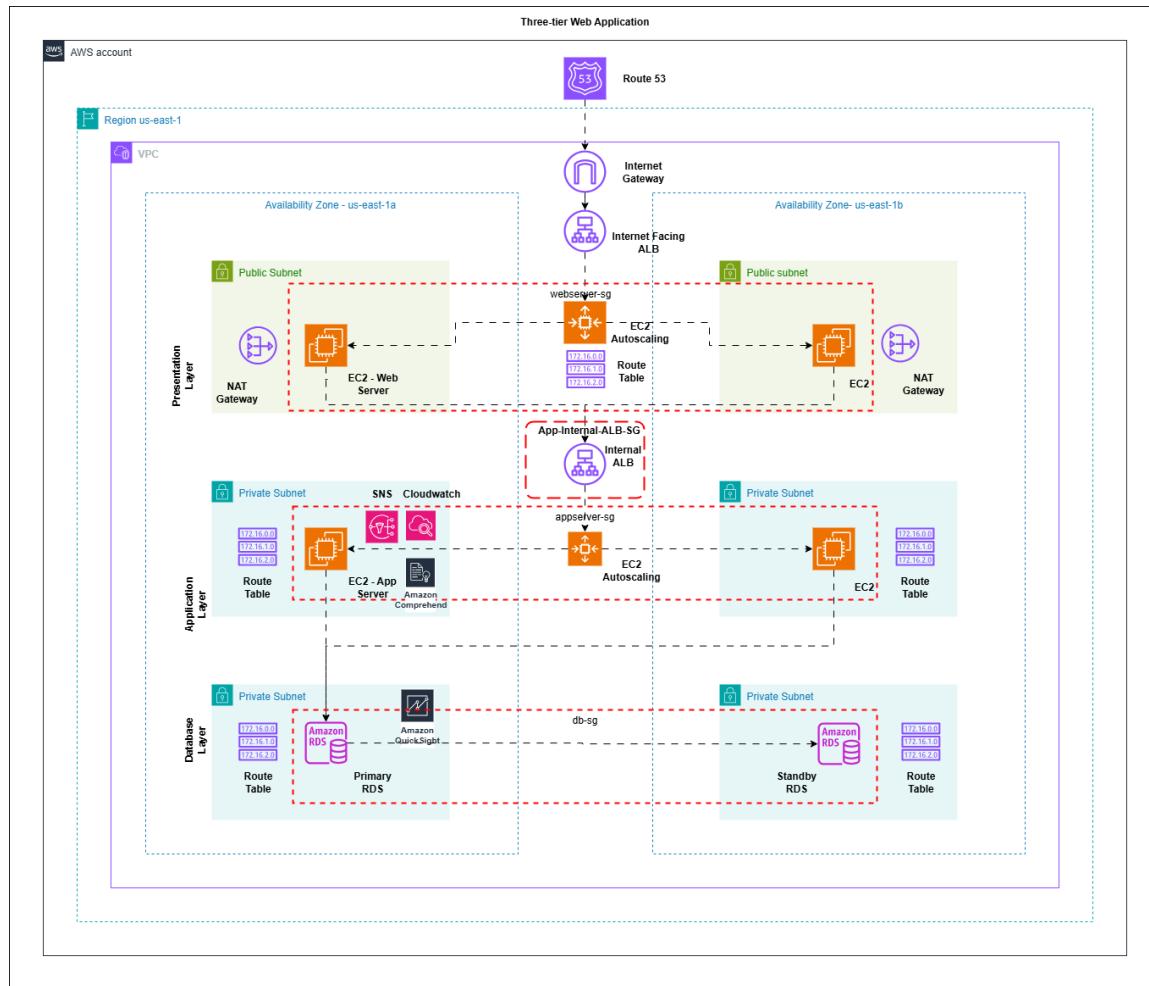
- **Enable Monitoring & Alerts:** Use Amazon CloudWatch to track key performance metrics and trigger alerts on anomalies.
- **Demonstrate Scalability:** Perform load testing to validate the Auto Scaling setup in real-world-like conditions.

AWS Architecture Diagram

Group Diagram



Individual Diagram



The architecture of the AWS Feedback Management System follows a secure, scalable, and highly available three-tier design:

- Amazon Route 53 – Handles domain name resolution for the feedback system's public endpoint and directs traffic to the Public ALB.
- Internet Gateway (IGW) – Provides internet connectivity to public subnets for the web tier and ALB.
- Public Application Load Balancer (ALB) – Distributes incoming traffic across web-tier EC2 instances in multiple Availability Zones.
- Web Tier (Public Subnets) –
 1. Auto Scaling Group (WebServer-ASG) runs EC2 instances hosting the feedback form frontend.
 2. Security Group (webserver-sg) allows HTTP/HTTPS traffic from anywhere and SSH access only from admin IPs.

- Internal Application Load Balancer (Internal ALB) –
 1. Routes requests from the web tier to the app tier.
 2. Security Group (App-Internal-ALB-SG) allows inbound HTTP only from webserver-sg.
- App Tier (Private Subnets) –
 1. Auto Scaling Group (AppServerASG) runs EC2 instances hosting the backend logic that processes feedback, calls AWS Comprehend, and writes to RDS.
 2. Security Group (appserverssg) allows HTTP only from the Internal ALB SG.
- Database Tier (Private Subnets) –
 1. Amazon RDS MySQL for storing data and analysis data.
 2. Security Group (dbsg) allows MySQL access only from appserverssg.
- NAT Gateways – Allow private-tier EC2 instances (app and DB layers) to reach the internet for OS updates and package installations without public exposure.
- Route Tables – Separate public and private routes, directing internet traffic via the IGW and private subnet traffic via NAT Gateways.

AWS Resource

The AWS Feedback Management System is deployed using the following AWS services and components:

- Amazon VPC – Provides an isolated virtual network for hosting all application components securely.
- Public and Private Subnets – Public subnets host web-tier EC2 instances and load balancers, while private subnets host the application-tier EC2 and RDS.
- Internet Gateway (IGW) – Enables public internet access for frontend servers and public ALB.
- NAT Gateways – Allow private subnets to access the internet for system updates without exposing them.
- Route Tables – Manage traffic flow between subnets, IGW, and NAT Gateways.
- Security Groups – Restrict inbound and outbound traffic for each component (web, app, DB, ALBs) using least-privilege principles.
- Public Application Load Balancer (ALB) – Routes external HTTP traffic to healthy web-tier EC2 instances.

- Internal Application Load Balancer (Internal ALB) – Routes traffic from the web tier to the app tier securely.
- Target Groups – Manage EC2 instance registration for load balancers and perform health checks.
- EC2 Auto Scaling Groups – Dynamically scale EC2 instances for web and app tiers based on CPU utilization and request load.
- Amazon EC2 – Web tier serves the frontend feedback form, and the app tier processes submissions, calls AWS Comprehend, and stores data in RDS.
- Amazon RDS (MySQL) – Managed database storing feedback submissions and sentiment results.
- AWS Comprehend – Analyzes feedback text to determine sentiment (Positive, Negative, Neutral, Mixed).
- Amazon QuickSight – Visualizes sentiment analysis and feedback trends through interactive dashboards.
- Amazon CloudWatch – Monitors application performance, collects logs, and sends alerts based on configured alarms.
- Amazon SNS – Sends notifications when Auto Scaling events occur, such as launching or terminating EC2 instances.

Implementation Steps

Networking Setup

- In the AWS Management Console, open the VPC service and create a VPC with CIDR block 10.0.0.0/16. Enable DNS hostnames and DNS resolution.
- Create 6 subnets:
- Create and attach an Internet Gateway (IGW) to the VPC to allow public internet access for public subnets.
- Create NAT Gateways in each public subnet to provide outbound internet access for private subnets without exposing them publicly.
- Create Route Tables

VPC dashboard < VPC > Your VPCs > vpc-018d63db50fc89933

vpc-018d63db50fc89933 / FMSvpc-vpc

Details **Info**

VPC ID	vpc-018d63db50fc89933	State	Available	Block Public Access	Off	DNS hostnames	Enabled
DNS resolution	Enabled	Tenancy	default	DHCP option set	dopt-0x434597d2c78b5bb	Main route table	rtb-010a8958e9cb76a4b
Main network ACL	act-0003990a575d40b4	Default VPC	No	IPv4 CDR	10.0.0.0/16	IPv6 pool	-
IPv6 CDR (Network border group)	-	Network Address Usage metrics	Disabled	Route 53 Resolver DNS Firewall rule groups	-	Owner ID	265735386481

Resource map **CDRs** **Flow logs** **Tags** **Integrations**

Resource map **Info**

```

graph LR
    VM[VM] --- S6[Subnets (6)]
    VM --- RT6[Route tables (6)]
    VM --- NC3[Network connections (3)]
    S6 --- RT6
    S6 --- NC3
    RT6 --- NC3
  
```

Subnets (6) Subnets within this VPC

- us-east-1a
 - FMSvpc-subnet-public1-us-east-1a
 - FMSvpc-subnet-private1-us-east-1a
 - FMSvpc-subnet-private3-us-east-1a
- us-east-1b
 - FMSvpc-subnet-public2-us-east-1b
 - FMSvpc-subnet-private2-us-east-1b

Route tables (6) Route network traffic to resources

- us-east-1a
 - FMSvpc-rtb-private2-us-east-1b
 - rtb-010a8958e9cb76a4b
 - FMSvpc-rtb-private3-us-east-1a
 - FMSvpc-rtb-private1-us-east-1a
- us-east-1b
 - FMSvpc-rtb-public
 - FMSvpc-rtb-private4-us-east-1b

Network connections (3) Connections to other networks

- FMSvpc-lgw
 - FMSvpc-nat-public1-us-east-1a
 - FMSvpc-nat-public2-us-east-1b

VPC dashboard < VPC > Subnets

Subnets (12) Info

Name	Subnet ID	State	VPC	Block Public...	IPv4 CDR	IPv6 CDR
FMSvpc-subnet-public1-us-east-1a	subnet-06e6f154974b9c53	Available	vpc-018d63db50fc89933 FMS...	Off	10.0.0.0/20	-
FMSvpc-subnet-private1-us-east-1a	subnet-06a710dc51fe81ac	Available	vpc-018d63db50fc89933 FMS...	Off	10.0.128.0/20	-
FMSvpc-subnet-private4-us-east-1b	subnet-08c0f40aae81d0461	Available	vpc-018d63db50fc89933 FMS...	Off	10.0.176.0/20	-
-	subnet-04f6b4e5cb1a9ed	Available	vpc-0cc62bd2f2bbe5258	Off	172.31.0.0/20	-
-	subnet-0bfb0d1f162f2ca4	Available	vpc-0cc62bd2f2bbe5258	Off	172.31.80.0/20	-
-	subnet-036969318014c0af6	Available	vpc-0cc62bd2f2bbe5258	Off	172.31.64.0/20	-
-	subnet-064294b077864cfa59	Available	vpc-0cc62bd2f2bbe5258	Off	172.31.32.0/20	-
FMSvpc-subnet-private3-us-east-1a	subnet-07fb128cabb1f173	Available	vpc-018d63db50fc89933 FMS...	Off	10.0.160.0/20	-
FMSvpc-subnet-private2-us-east-1b	subnet-0ca6f96545d2206b4	Available	vpc-018d63db50fc89933 FMS...	Off	10.0.144.0/20	-
FMSvpc-subnet-public2-us-east-1b	subnet-0407e2f2f51598fe	Available	vpc-018d63db50fc89933 FMS...	Off	10.0.16.0/20	-
-	subnet-0773f26380efffb15	Available	vpc-0cc62bd2f2bbe5258	Off	172.31.48.0/20	-
-	subnet-0a38f6d46cf9ce15	Available	vpc-0cc62bd2f2bbe5258	Off	172.31.16.0/20	-

Select a subnet

The screenshot shows the AWS VPC dashboard for the 'igw-0862874e469ec35ab' Internet Gateway. The left sidebar includes sections for EC2 Global View, Virtual private cloud (with 'Your VPCs' selected), Internet gateways (selected), Security, PrivateLink and Lattice, and CloudShell/Feedback. The main content area displays the Internet gateway details, showing it is attached to a VPC (VPC ID: vpc-018d63db50fc89933 | FMSvpc-igw). It also lists a tag named 'Name' with the value 'FMSvpc-igw'. The top right corner shows the account ID (2657-5158-4481) and location (United States (N. Virginia)).

The screenshot shows the AWS VPC dashboard with the 'NAT gateways' section selected. The table lists two NAT gateways:

Name	NAT gateway ID	Connectivity...	State	State message	Primary public I...	Primary private I...	Primary network...	VPC
FMVpc-nat-public1...	nat-0adeef74716374ed5	Public	Available	-	44.208.61.82	10.0.6.174	eni-01bc5615925e0...	vpc-018d65db50f899...
FMVpc-nat-public2...	nat-0579ae063adff674	Public	Available	-	52.206.148.192	10.0.27.36	eni-0ab0cc5d5f57...	vpc-018d65db50f899...

Below the table, a section titled 'Select a NAT gateway' is visible.

Screenshot of the AWS VPC NAT gateway details page for nat-Oadee7471637f4ed5.

Details

NAT gateway ID nat-Oadee7471637f4ed5	Connectivity type Public	State Available	State message Info
NAT gateway ARN arn:aws:ec2:us-east-1:265753586481:natgateway/nat-Oadee7471637f4ed5	Primary public IPv4 address 44.208.61.82	Primary private IPv4 address 10.0.6.174	Primary network interface ID eni-01bc3613923e0bd7
VPC vpc-018d63db50fc89933 / FMSvpc-vpc	Subnet subnet-06e6fc154974b9cc3 / FMSvpc-subnet-public1-us-east-1a	Created Monday, July 28, 2025 at 10:15:35 EDT	Deleted -

Secondary IPv4 addresses

Private IPv4 address	Allocation ID	Association ID	Public IPv4 address	Network interface ID
Secondary IPv4 addresses are not available for this nat gateway.				

Actions: Edit secondary IPv4 address associations

CloudShell Feedback

Screenshot of the AWS VPC NAT gateway details page for nat-0578ae0863edf67d4.

Details

NAT gateway ID nat-0578ae0863edf67d4	Connectivity type Public	State Available	State message Info
NAT gateway ARN arn:aws:ec2:us-east-1:265753586481:natgateway/nat-0578ae0863edf67d4	Primary public IPv4 address 52.206.148.192	Primary private IPv4 address 10.0.27.36	Primary network interface ID eni-0eab0c03d9574f91
VPC vpc-018d63db50fc89933 / FMSvpc-vpc	Subnet subnet-0407e2f25f1598fe / FMSvpc-subnet-public2-us-east-1b	Created Monday, July 28, 2025 at 10:15:35 EDT	Deleted -

Secondary IPv4 addresses

Private IPv4 address	Allocation ID	Association ID	Public IPv4 address	Network interface ID
Secondary IPv4 addresses are not available for this nat gateway.				

Actions: Edit secondary IPv4 address associations

CloudShell Feedback

Name	Route table ID	Explicit subnet associations	Edge associations	Main	VPC	Owner ID
FMSvpc-rtb-private2-us-east-1b	rtb-05213f94d21b0381	subnet-0ca6f965d3d220...	-	No	vpc-018d63db50fc89933 FMS...	265735586481
-	rtb-010a8958e9c76ea8b	-	-	Yes	vpc-018d63db50fc89933 FMS...	265735586481
-	rtb-0f5598ce0994bddd	-	-	Yes	vpc-0cc62bd2f22be5258	265735586481
FMSvpc-rtb-private3-us-east-1a	rtb-0e541a6e80be10abf	subnet-07fb128c3abcf1c...	-	No	vpc-018d63db50fc89933 FMS...	265735586481
FMSvpc-rtb-private1-us-east-1a	rtb-0e740fa8bde621f7	subnet-06a7100c51fe8...	-	No	vpc-018d63db50fc89933 FMS...	265735586481
FMSvpc-rtb-public	rtb-01fbbed55213560a1	2 subnets	-	No	vpc-018d63db50fc89933 FMS...	265735586481
FMSvpc-rtb-private4-us-east-1b	rtb-0e02e0d0d4de9c8a3	subnet-08ccff40aae81d8...	-	No	vpc-018d63db50fc89933 FMS...	265735586481

Security Configuration

Security Groups

- Create WebServerSG:
 - Inbound: Allow HTTP (80), SSH (22) and HTTPS (443) from 0.0.0.0/0.
 - Outbound: Allow all traffic (default).
- Create AppServerSG:
 - Inbound: Allow SSH (22) and Custom TCP(5000) from WebServerSG.
 - Outbound: Allow all traffic (default).
- Create DBSG, DbAnalysisSG:
 - Inbound: Allow MySQL (3306) from AppServerSG.
 - Outbound: Allow all traffic (default).
- Create ALB security group:
 - Internal ALB-SG: Allow HTTP(80) from WebServerSG only.

VPC dashboard < VPC > Security Groups

Security Groups (8) Info

Name	Security group ID	Security group name	VPC ID	Description	Owner
sg-08c2172576fbfb956	AppServerSG	vpc-018d63db50f89933	AppServerSG	AppServerSG	265753586481
sg-0a9650af60955e4b	App-Internal-ALB-SG	vpc-018d63db50f89933	App-Internal-ALB-SG	App-Internal-ALB-SG	265753586481
sg-02229c87244e2ac78	WebServerSG	vpc-018d63db50f89933	WebServerSG	WebServerSG	265753586481
sg-0beac234a9589330	QuickSightAccessSG	vpc-018d63db50f89933	QuickSightAccessSG	QuickSightAccessSG	265753586481
sg-0ba45386b794fbca	default	vpc-018d63db50f89933	default VPC security group	default VPC security group	265753586481
sg-060ad3d22e174b4de	default	vpc-0cc62bd2f2be5258	default VPC security group	default VPC security group	265753586481
sg-0e3782dff4e51efb7	DbAnalysisSG	vpc-018d63db50f89933	DbAnalysisSG	DbAnalysisSG	265753586481
sg-0edfb2ba833c86e93	DbSG	vpc-018d63db50f89933	DbSG	DbSG	265753586481

Select a security group

© 2025, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences 10:33 AM 13/06/2025

Console Home | Console < Instances | EC2 | us-east-1 < SecurityGroup | EC2 | us- < Target group details | EC2 < Auto Scaling group details | Subscriptions 77b15a12 < analysis analysis < ALARM: "Warning-AGG-1" < + All Bookmarks

EC2 > Security Groups > sg-02229c87244e2ac78 - WebServerSG Actions

Details

Security group name	WebServerSG	Security group ID	sg-02229c87244e2ac78	Description	vpc-018d63db50f89933
Owner	265753586481	Inbound rules count	3 Permission entries	Outbound rules count	1 Permission entry

Inbound rules | Outbound rules | Sharing - new | VPC associations - new | Tags

Inbound rules (3)

Name	Security group rule ID	IP version	Type	Protocol	Port range	Source	Description
-	sgr-000bdfa138cab2148	IPv4	SSH	TCP	22	0.0.0.0/0	-
-	sgr-0791093af905846	IPv4	HTTP	TCP	80	0.0.0.0/0	-
-	sgr-0f731526d42b28879	IPv4	HTTPS	TCP	443	0.0.0.0/0	-

Manage tags | Edit inbound rules

CloudShell Feedback

© 2025, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences 8:33 PM 13/06/2025

Screenshot of the AWS EC2 Security Groups console showing the details for the security group **sg-08c217257b62fb956 - AppServerSG**.

Details

Security group name	AppServerSG	Security group ID	sg-08c217257b62fb956
Owner	265753586481	Description	AppServerSG
		VPC ID	vpc-018d63db50f89933

Inbound rules count: 2 Permission entries
Outbound rules count: 1 Permission entry

Inbound rules (2)

Name	Security group rule ID	IP version	Type	Protocol	Port range	Source	Description
-	sgr-0ce4c66af6acc6a2b	-	SSH	TCP	22	sg-0229b87244e2a78...	-
-	sgr-05f4829fb7b545316	-	Custom TCP	TCP	5000	sg-0d650af609f55eb...	-

Actions: Manage tags, Edit inbound rules

CloudShell Feedback

Screenshot of the AWS EC2 Security Groups console showing the details for the security group **sg-0edfb2ba833c86a93 - DbSG**.

Details

Security group name	DbSG	Security group ID	sg-0edfb2ba833c86a93
Owner	265753586481	Description	DbSG
		VPC ID	vpc-018d63db50f89933

Inbound rules count: 1 Permission entry
Outbound rules count: 1 Permission entry

Inbound rules (1)

Name	Security group rule ID	IP version	Type	Protocol	Port range	Source	Description
-	sgr-034be7646326c9261	-	MySQL/Aurora	TCP	3306	sg-08c217257b62fb95...	-

Actions: Manage tags, Edit inbound rules

CloudShell Feedback

sg-0e3782dff4e51efb7 - DbAnalysisSG

Details

Security group name	sg-0e3782dff4e51efb7	Security group ID	sg-0e3782dff4e51efb7	Description	VPC ID
Owner	265753586481	Inbound rules count	1 Permission entry	Outbound rules count	1 Permission entry

Inbound rules | Outbound rules | Sharing - new | VPC associations - new | Tags

Inbound rules (1)

Name	Security group rule ID	IP version	Type	Protocol	Port range	Source	Description
-	sgr-0df90a05e640397ea	-	MySQL/Aurora	TCP	3306	sg-08c217257b62fb95...	-

sg-0a9650af609f55e4b - App-Internal-ALB-SG

Details

Security group name	sg-0a9650af609f55e4b	Security group ID	sg-0a9650af609f55e4b	Description	VPC ID
Owner	265753586481	Inbound rules count	1 Permission entry	Outbound rules count	1 Permission entry

Inbound rules | Outbound rules | Sharing - new | VPC associations - new | Tags

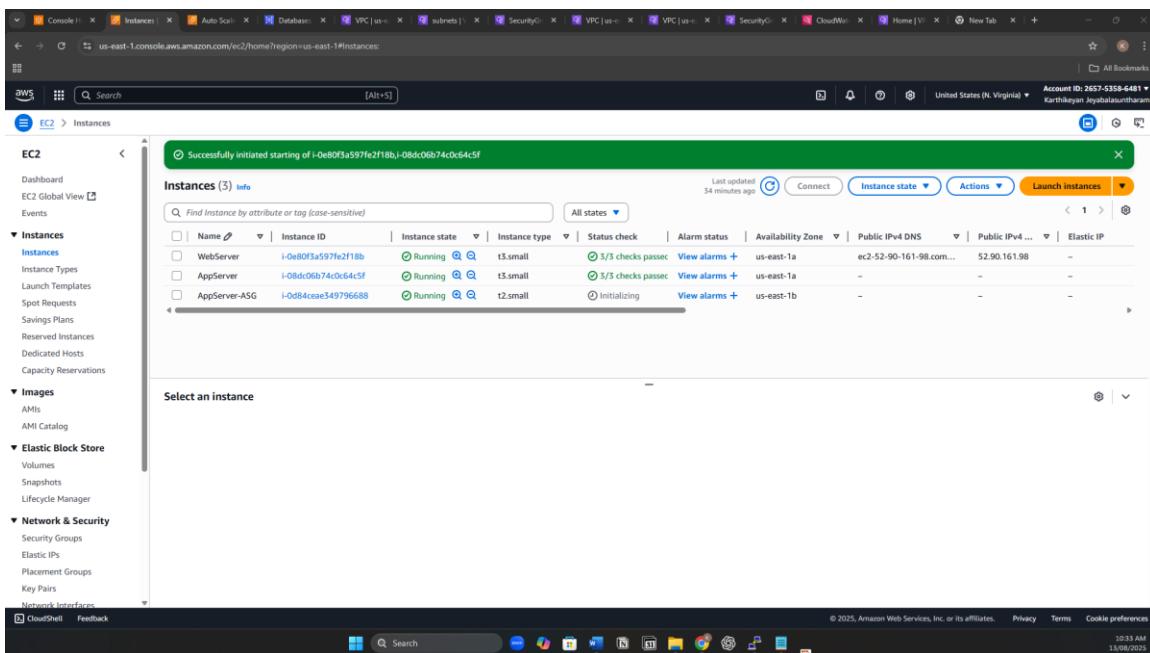
Inbound rules (1)

Name	Security group rule ID	IP version	Type	Protocol	Port range	Source	Description
-	sgr-046241fe1464a7fde	-	HTTP	TCP	80	sg-0222987244e2e78...	-

Web & App Deployment

- Launch EC2 instances in the Web Tier (public subnets) using Amazon Linux 2023.
- Assign WebServer-SG.
- Install Apache web server and configure it to serve the frontend feedback form.

- Launch EC2 instances in the App Tier (private subnets) using Amazon Linux 2023. Assign AppServer-SG.
- Install Python3, Flask, MySQL-Connector, and Boto3 for AWS Comprehend integration.
- Deploy the feedback application code to /home/ec2-user/feedback-app.
- Create feedback-app.service systemd unit file to start the app automatically on boot
- Enable and start the service:



Screenshot of the AWS EC2 Instances page showing the instance summary for i-0e80f3a597fe2f18b. The instance is running and has a public IPv4 address of 52.90.161.98.

Instance summary for i-0e80f3a597fe2f18b (WebServer) Info

Updated less than a minute ago

Instance ID i-0e80f3a597fe2f18b	Public IPv4 address 52.90.161.98 [open address]
IPv6 address -	Instance state Running
Hostname type IP name: ip-10-0-4-230.ec2.internal	Private IP DNS name (IPv4 only) ip-10-0-4-230.ec2.internal
Answer private resource DNS name -	Instance type t3.small
Auto-assigned IP address 52.90.161.98 [Public IP]	VPC ID vpc-018d63db50fc8933 (FMSvpc-vpc)
IAM Role -	Subnet ID subnet-06e6fc154974b9c3 (FMSvpc-subnet-public1-us-east-1a)
IMDSv2 Required	Instance ARN arn:aws:ec2:us-east-1:265753586481:instance/i-0e80f3a597fe2f18b
Operator -	

Details **Status and alarms** **Monitoring** **Security** **Networking** **Storage** **Tags**

Instance details Info

AMI ID ami-08a6efd148b1f7504	Monitoring disabled
AMI name al2023-ami-2023.8.20250721.2-kernel-6.1-x86_64	Platform details Linux/UNIX

© 2025, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences 10:34 AM 13/06/2025

Screenshot of the AWS EC2 Instances page showing the instance summary for i-08dc06b74c0c64c5f. The instance is running and has a public IPv4 address of 10.0.137.70.

Instance summary for i-08dc06b74c0c64c5f (AppServer) Info

Updated less than a minute ago

Instance ID i-08dc06b74c0c64c5f	Public IPv4 address -
IPv6 address -	Instance state Running
Hostname type IP name: ip-10-0-137-70.ec2.internal	Private IP DNS name (IPv4 only) ip-10-0-137-70.ec2.internal
Answer private resource DNS name -	Instance type t3.small
Auto-assigned IP address -	VPC ID vpc-018d63db50fc8933 (FMSvpc-vpc)
IAM Role EC2ComprehendRole	Subnet ID subnet-06a710d0c51fe81ac (FMSvpc-subnet-private1-us-east-1a)
IMDSv2 Required	Instance ARN arn:aws:ec2:us-east-1:265753586481:instance/i-08dc06b74c0c64c5f
Operator -	

Details **Status and alarms** **Monitoring** **Security** **Networking** **Storage** **Tags**

Instance details Info

AMI ID ami-08a6efd148b1f7504	Monitoring disabled
AMI name al2023-ami-2023.8.20250721.2-kernel-6.1-x86_64	Platform details Linux/UNIX
	Termination protection Disabled

© 2025, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences 10:34 AM 13/06/2025

```

ec2-user@ip-10-0-137-70:~/feedback-app
GNU nano 8.3
#com flask import Flask, request, jsonify
import boto3
import pymysql

app = Flask(__name__)

# ----- RDS DB1 (Raw Feedback) -----
DB1_HOST = 'fmsdb.cg2uiquqgc.us-east-1.rds.amazonaws.com'
DB1_USER = 'admin'
DB1_PASS = 'DB1@123'
DB1_NAME = 'feedback_raw'           # same DB name as your friend
DB1_TABLE = 'feedbacks'            # same table name, new columns

def get_db1_connection():
    return pymysql.connect(
        host=DB1_HOST,
        user=DB1_USER,
        password=DB1_PASS,
        database=DB1_NAME,
        cursorclass=pymysql.cursors.DictCursor
    )

# ----- RDS DB2 (Analyzed Feedback) -----
DB2_HOST = 'feedbackanalysisdb.cg2uiquqgc.us-east-1.rds.amazonaws.com'
DB2_USER = 'admin'
DB2_PASS = 'DB1@123'
DB2_NAME = 'feedback_analysis'      # same DB name as your friend
DB2_TABLE = 'analysis'             # same table name, new columns

def get_db2_connection():
    return pymysql.connect(
        host=DB2_HOST,
        user=DB2_USER,
        password=DB2_PASS,
        database=DB2_NAME,
        cursorclass=pymysql.cursors.DictCursor
    )

# ----- Amazon Comprehend -----
# keep region consistent with your resources
comprehend = boto3.client('comprehend', region_name='us-east-2')

@app.route('/health', methods=['Get'])
def health():
    return jsonify({'status': 'ok'}), 200

```

Load Balancing

- Create a Public Application Load Balancer for Web Tier.
- Create an Internal Application Load Balancer App Tier.
- Register target groups.

The screenshot shows the AWS CloudWatch Metrics interface with a graph titled 'CPU Utilization' for two targets: 'WebServer1' and 'WebServer2'. The Y-axis represents CPU Utilization from 0% to 100%, and the X-axis represents time from 10:30 AM to 11:30 AM. Both targets show fluctuating utilization levels, with peaks around 100% and troughs around 0%.

Internal-App-ALB

Listeners and rules (1) Info

ProtocolPort	Default action	Rules	ARN	Security policy	Default SSL/TLS certificate	mTLS	Trust store
HTTP:80	Forward to target group app-tg-5000 (100%) Target group stickiness: Off	1 rule	ARN	Not applicable	Not applicable	Not applicable	Not applicable

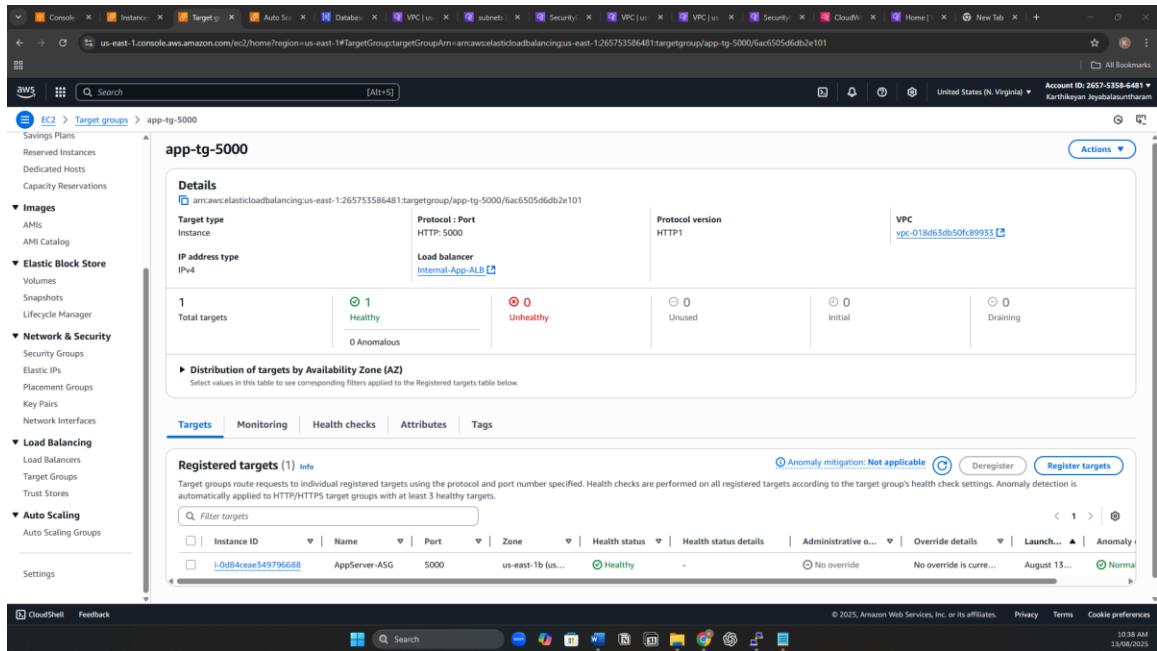
WebServerTG

Targets

Total targets	Healthy	Unhealthy	Unused	Initial	Draining
1	1	0	0	0	0

Registered targets (1) Info

Instance ID	Name	Port	Zone	Health status	Health status details	Administrative o...	Override details	Launch...	Anomaly
i-0e00f5a597fe2f1Bb	WebServer	80	us-east-1a (us...)	Healthy	-	No override	No override is curre...	August 13...	Normal



Auto Scaling

- Create Launch Templates for both Web Tier and App Tier:
 - Specify AMI, instance type, key pair, security group, and user data to install packages and deploy app.
 - Create Auto Scaling Groups:
 - WebServer-ASG: Public subnets, attach Web Tier target group.
 - AppServer-ASG: Private app subnets, attach App Tier target group.
 - Configure scaling policies:
 - Target tracking based on average CPU utilization.
 - Enable notifications:
 - Create an Amazon SNS topic for Auto Scaling events.
 - Subscribe with email.
 - Attach SNS topic to the ASG notifications for instance launch/terminate events.

EC2 > Target groups > app-tg-5000

app-tg-5000

Details
 Target type: Instance
 Protocol: HTTP: 5000
 IP address type: IPv4
 Load balancer: Internal-App-ALB

Total targets	Healthy	Unhealthy	Unused	Initial	Draining
4	1	0	0	0	3
0 Anomalous					

Distribution of targets by Availability Zone (AZ)
 Select values in this table to see corresponding filters applied to the Registered targets table below.

Targets | Monitoring | Health checks | Attributes | Tags

Registered targets (4) Info
 Target groups route requests to individual registered targets using the protocol and port number specified. Health checks are performed on all registered targets according to the target group's health check settings. Anomaly detection is automatically applied to HTTP/HTTPS target groups with at least 3 healthy targets.

Instance ID	Name	Port	Zone	Health status	Health status details	Administrative o...	Override details	Launch time
I-02952614ab554eaf5	AppServer-ASG	5000	us-east-1a (un...	Draining	Target deregistration is in progress	No override	No override	August 13, 2025, 18:55 (U)
I-0e4bc564f2e5036f	AppServer-ASG	5000	us-east-1b (un...	Healthy	-	No override	No override	August 13, 2025, 18:55 (U)
I-015b04a6545518056	WebServer	80	us-east-1a (un...	Healthy	-	No override	No override is current	August 13, 2025, 18:55 (U)
I-0e0ff5a597fe2f1bb	WebServer	80	us-east-1a (un...	Healthy	-	No override	No override is current	August 13, 2025, 18:55 (U)

© 2025, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences 8:17 PM 13/08/2025

EC2 > Target groups > WebServerTG

WebServerTG

Details
 Target type: Instance
 Protocol: HTTP: 80
 IP address type: IPv4
 Load balancer: FrontendLB

Total targets	Healthy	Unhealthy	Unused	Initial	Draining
2	2	0	0	0	0
0 Anomalous					

Distribution of targets by Availability Zone (AZ)
 Select values in this table to see corresponding filters applied to the Registered targets table below.

Targets | Monitoring | Health checks | Attributes | Tags

Registered targets (2) Info
 Target groups route requests to individual registered targets using the protocol and port number specified. Health checks are performed on all registered targets according to the target group's health check settings. Anomaly detection is automatically applied to HTTP/HTTPS target groups with at least 3 healthy targets.

Instance ID	Name	Port	Zone	Health status	Health status details	Administrative o...	Override details	Launch...	Anomaly
I-015b04a6545518056	WebServer	80	us-east-1a (un...	Healthy	-	No override	No override is current	August 13, 2025, 18:55 (U)	Normal
I-0e0ff5a597fe2f1bb	WebServer	80	us-east-1a (un...	Healthy	-	No override	No override is current	August 13, 2025, 18:55 (U)	Normal

© 2025, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences 8:17 PM 13/08/2025

Screenshot of the AWS EC2 Auto Scaling group details page for 'app-asg'.

app-asg Capacity overview

Desired capacity	Scaling limits (Min - Max)	Desired capacity type
1	1 - 4	Units (number of instances)

Date created: Sun Aug 10 2025 11:07:14 GMT-0400 (Eastern Daylight Time)

Launch template

Launch template	AMI ID	Instance type	Owner
lt-00625ce9e96b47d25 app-asg-it	ami-00f123071fb25c000	t2.small	arn:aws:iam::265753586481:root

Version: Default

Description: v1 - ami v1

Network

Availability Zones	Subnet ID	Availability Zone distribution
--------------------	-----------	--------------------------------

Details **Integrations - new** **Automatic scaling** **Instance management** **Instance refresh** **Activity** **Monitoring**

© 2025, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences 10:39 AM 11/06/2025

Screenshot of the AWS EC2 Auto Scaling group scaling policies page for 'app-asg'.

Dynamic scaling policy created or edited successfully.

scaling policies along with dynamic scaling policies in the following situations: when your application demand changes quickly, but with a recurring pattern, or when your EC2 instances require more time to initialize.

Dynamic scaling policies (1) Info

Target Tracking Policy	Actions Create dynamic scaling policy
-------------------------------	---

Target Tracking Policy

Policy type: Target tracking scaling

Enabled or disabled: Enabled

Execute policy when: As required to maintain Average CPU utilization at 2

Take the action: Add or remove capacity units as required

Instances need: 60 seconds to warm up before including in metric

Scale in: Enabled

Predictive scaling policies (0) Info

Evaluation period	Actions Create predictive scaling policy
--------------------------	--

Evaluation based on 2 days

© 2025, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences 10:40 AM 11/06/2025

Capacity overview

Desired capacity	Scaling limits (Min - Max)	Desired capacity type	Status
1	1 - 4	Units (number of instances)	-

Launch template

Launch template	AMI ID	Instance type	Owner
lt-0211440a14d9e70b5 web-asg-lt	ami-011a6bae2d2dd5af0	t2.small	arn:aws:iam::265753586481:root

Network

Availability Zones	Subnet IDs	Availability Zone distribution
-	-	-

Automatic scaling

Scaling policies resize your Auto Scaling group to meet changes in demand. With reactive dynamic scaling policies, you can track specific CloudWatch metrics and take action when the CloudWatch alarm threshold is met. Use predictive scaling policies along with dynamic scaling policies in the following situations: when your application demand changes quickly, but with a recurring pattern, or when your EC2 instances require more time to initialize.

Dynamic scaling policies (1) Info

Policy type	Actions	Create dynamic scaling policy
Target tracking scaling		

ALB Request Count per Target

Policy type	Actions	Create dynamic scaling policy
Target tracking scaling		

Enabled or disabled

Enabled

Execute policy when

As required to maintain Application Load Balancer request count per target at 60

Take the action

Add or remove capacity units as required

Instances need

60 seconds to warm up before including in metric

Scale in

Enabled

Predictive scaling policies (0) Info

Screenshot of the AWS CloudWatch Alarms page showing the "Warning-ASG-MaxCapacityReached" alarm.

CloudWatch Alarms

Warning-ASG-MaxCapacityReached

Details

- Name:** Warning-ASG-MaxCapacityReached
- Type:** Metric alarm
- Description:** app-asg_DesiredCapacity >= 4 (Max). Auto Scaling Capacity Alert — app-asg. The Auto Scaling Group app-asg has reached its configured maximum capacity of 4. No further scale-out actions can occur, which may impact performance if demand continues to rise. Recommended actions:
 - Review current load and CPU utilization
 - Check application performance and latency
 - Consider increasing MaxCapacity or optimizing workloads
- State:** OK
- Threshold:** GroupDesiredCapacity >= 4 for 1 datapoints within 1 minute
- Last state update:** 2025-08-14 00:10:01 (UTC)
- Actions:** Average
- Period:** 1 minute

Datapoints to alarm: 1 out of 1

Metric name: GroupDesiredCapacity

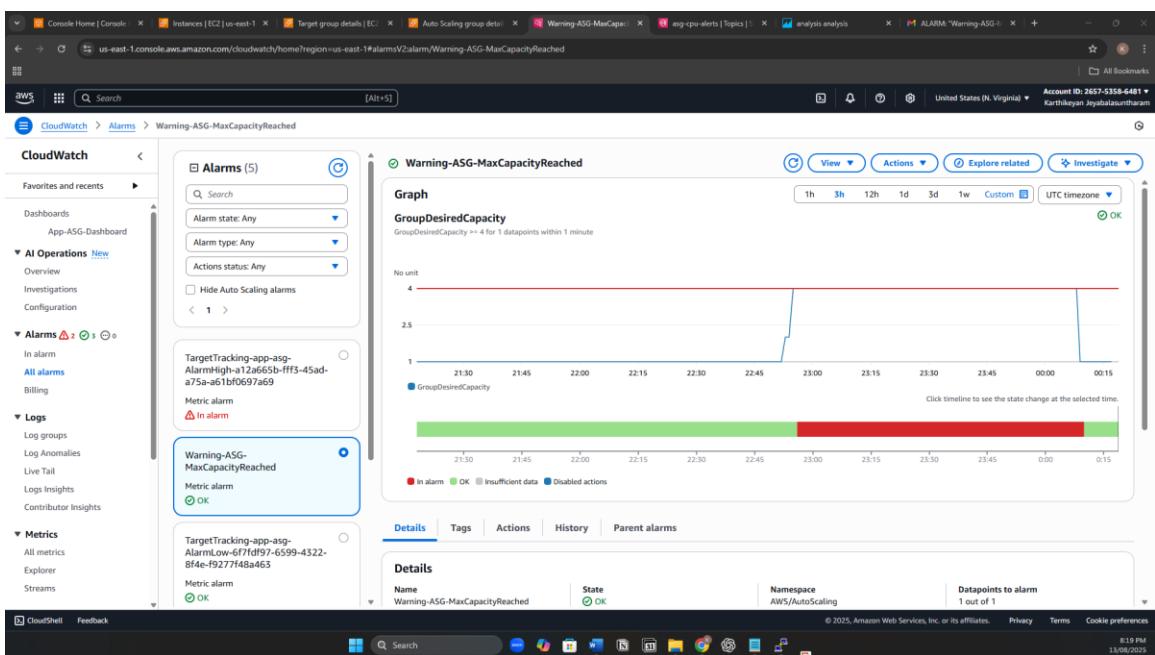
Namespace: AWS/AutoScaling

ARN: arn:aws:cloudwatch:us-east-1:265753586481:alarm:Warning-ASG-MaxCapacityReached

Missing data treatment: Treat missing data as missing

Percentiles with low samples: evaluate

View EventBridge rule



The screenshot shows the AWS SNS console with a subscription configuration. The subscription ARN is arn:aws:sns:us-east-1:265753586481:asg-cpu-alerts:77b15a12-b691-4aeb-8d34-71e0233a34da. The endpoint is set to an email address (karthikedujk@gmail.com). The topic is asg-cpu-alerts, and the subscription principal is arn:aws:iam::265753586481:root. The status is confirmed via EMAIL. There is no filter policy applied.

The screenshot shows the AWS SNS console with a topic configuration. The topic name is asg-cpu-alerts. It is a Standard topic with an ARN of arn:aws:sns:us-east-1:265753586481:asg-cpu-alerts. The display name is asg-cpu-alerts, and the topic owner is 265753586481. A single subscription is listed, which is the one shown in the previous screenshot.

Database Setup

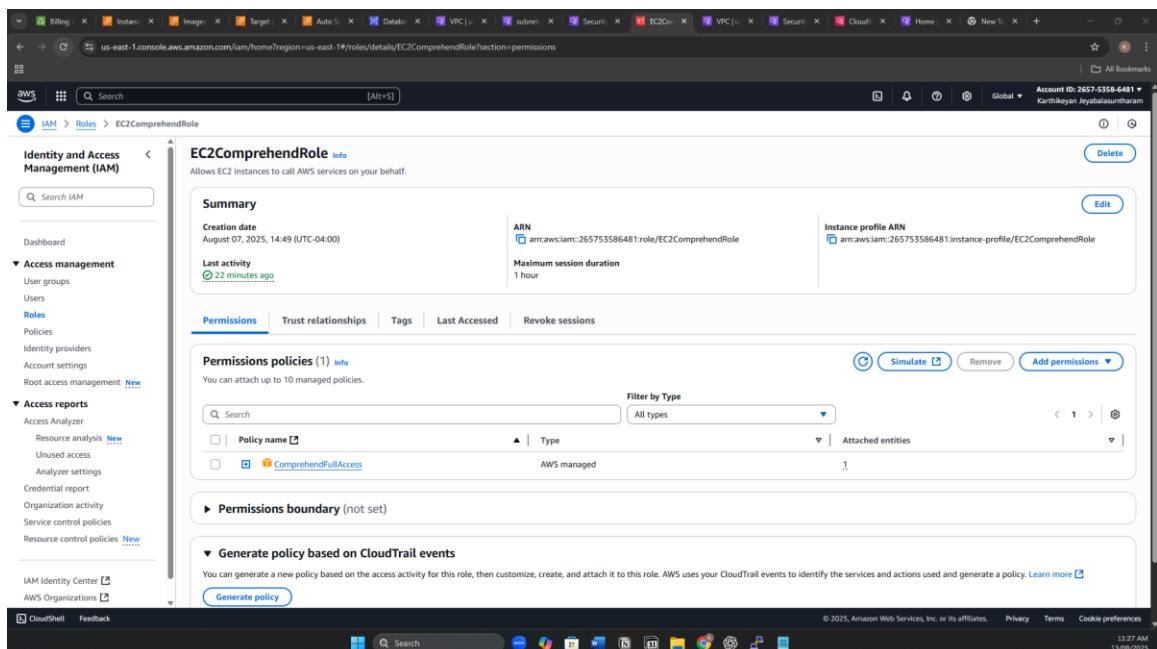
- Launch Amazon RDS MySQL in private DB subnets.
- Assign DB-SG to RDS instance.
- Set database credentials and note the endpoint.
- Update application configuration to connect to RDS endpoint with secure credentials.

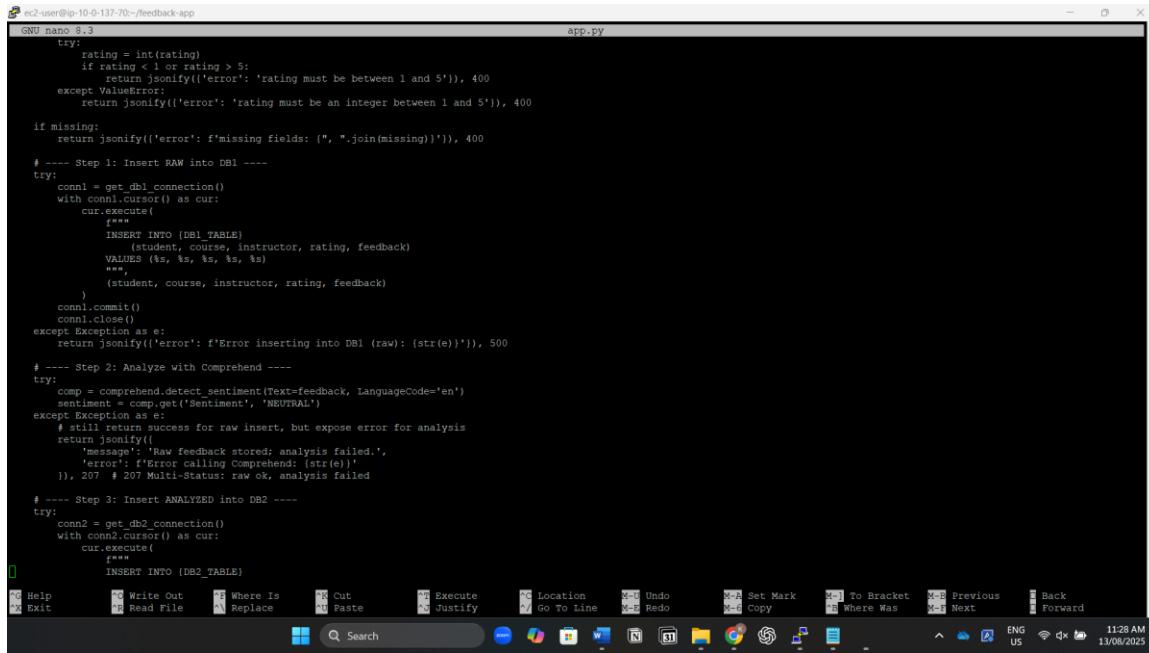
Screenshot of the AWS RDS console showing the details for the database 'fmldb'. The 'Summary' tab is selected, displaying basic information like DB identifier, status (Available), role (Instance), engine (MySQL Community), and region & AZ (us-east-1a). The 'Connectivity & security' tab is active, showing the endpoint (fmldb.cgp2uuiquqrg.us-east-1.rds.amazonaws.com), port (3306), VPC (RMSvc-vpc (vpc-018d63db50fc89933)), subnet group (fmssubnetgroup), and network type (IPv4). The 'Networking' section shows the availability zone (us-east-1a) and subnets (subnet-07fb128c5abcf1c73, subnet-08cf40aae81d8461). The 'Security' section includes VPC security groups (Db565 (sg-0e9ff2ba853c86a93)), certificate authority (Info), and instance certificate expiration date (July 28, 2026, 11:51 (UTC-04:00)).

Screenshot of the AWS RDS console showing the details for the database 'feedbackanalysistdb'. The 'Summary' tab is selected, displaying basic information like DB identifier, status (Available), role (Instance), engine (MySQL Community), and region & AZ (us-east-1a). The 'Connectivity & security' tab is active, showing the endpoint (feedbackanalysistdb.cgp2uuiquqrg.us-east-1.rds.amazonaws.com), port (3306), VPC (RMSvc-vpc (vpc-018d63db50fc89933)), subnet group (fmssubnetgroup), and network type (IPv4). The 'Networking' section shows the availability zone (us-east-1a) and subnets (subnet-07fb128c5abcf1c73, subnet-08cf40aae81d8461). The 'Security' section includes VPC security groups (DbAnalysisSG (sg-0e3782df4e51efb7)), certificate authority (Info), and instance certificate expiration date (August 07, 2026, 14:20 (UTC-04:00)).

AWS Comprehend Integration

- Create an IAM Role for EC2 with permission comprehend:FullAccess.
 - Attach the IAM role to App Tier EC2 instances.
 - Update application backend code to:
 - Call AWS Comprehend’s API with user feedback text.
 - Store sentiment result (Positive, Negative, Neutral) along with feedback in RDS.





```

e2-user@ip-10-0-137-70:~/feedback-app
GNU nano 8.3
app.py

try:
    rating = int(rating)
    if rating < 1 or rating > 5:
        return jsonify({'error': 'rating must be between 1 and 5'}), 400
except ValueError:
    return jsonify({'error': 'rating must be an integer between 1 and 5'}), 400

if missing:
    return jsonify({'error': f'missing fields: {", ".join(missing)}'}), 400

# ---- Step 1: Insert RAW into DB1 ----
try:
    conn1 = get_db1_connection()
    with conn1.cursor() as cur:
        cur.execute(
            f"""
            INSERT INTO [DB1_TABLE]
            (student, course, instructor, rating, feedback)
            VALUES (%s, %s, %s, %s, %s)
            """,
            (student, course, instructor, rating, feedback)
        )
    conn1.commit()
    conn1.close()
except Exception as e:
    return jsonify({'error': f'Error inserting into DB1 (raw): ({str(e)})'}), 500

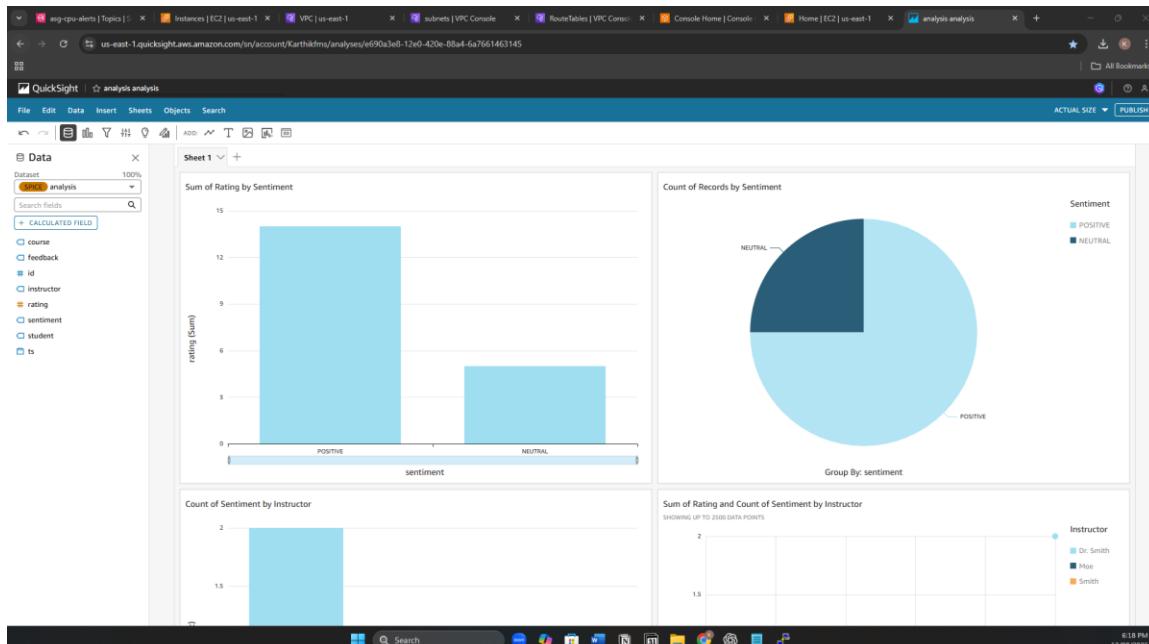
# ---- Step 2: Analyze with Comprehend ----
try:
    comp = comprehend_detect_sentiment(Text=feedback, LanguageCode='en')
    sentiment = comp.get('Sentiment', 'NEUTRAL')
except Exception as e:
    # still return success for raw insert, but expose error for analysis
    return jsonify({
        'message': 'Raw feedback stored; analysis failed.',
        'error': f'Error calling Comprehend: ({str(e)})'
    }), 207 # 207 Multi-Status: raw OK, analysis failed

# ---- Step 3: Insert ANALYZED into DB2 ----
try:
    conn2 = get_db2_connection()
    with conn2.cursor() as cur:
        cur.execute(
            f"""
            INSERT INTO [DB2_TABLE]
            """
        )

```

Visualization with QuickSight

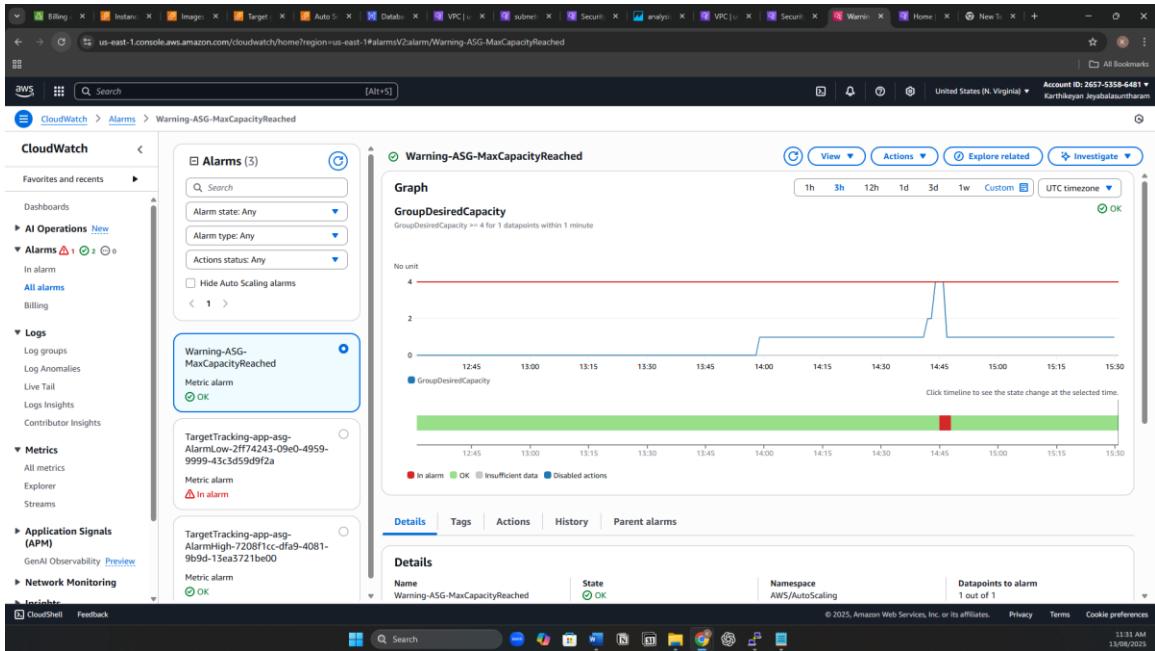
- Ensure QuickSight has a VPC connection to access RDS in private subnets.
- In QuickSight, create a new data source and connect to the RDS MySQL instance.
- Create datasets from feedback and sentiment tables.
- Build visual dashboards showing sentiment distribution, rating averages, and trends over time.



Monitoring & Logging

- Create CloudWatch Alarms for:
 - High CPU utilization on web/app instances.
- Configure alarm actions to send notifications to SNS topic and send email notification.
- Review metrics and alarms regularly to ensure system health and performance.

The screenshot displays two separate AWS CloudWatch dashboards. The top dashboard, titled 'CloudWatch Metrics', shows an 'Overview' section with a timeline from 13:00 to 15:00. It highlights an alarm named 'Warning-ASG-MaxCapacityReached' for the AutoScaling service, which triggered at least once. The bottom dashboard, titled 'CloudWatch Metrics Insights', features an 'Application Insights' section with a 'Get started with Application Insights' callout. Both dashboards include navigation menus on the left and standard browser UI elements like tabs, search bars, and status bars at the bottom.



```

ec2-user@ip-10-0-4-230:~$ Run "lsb_release -a" for full release and version update info
Amazon Linux 2023
https://aws.amazon.com/linux/amazon-linux-2023

Last login: Wed Aug 13 21:24:26 2025 from 10.0.4.230
[ec2-user@ip-10-0-137-70 ~]$ mysql -h feedbackanalysisdb.cgp2uuiquqcg.us-east-1.rds.amazonaws.com -u admin -p -e "USE feedback_analysis; SELECT id, course, instructor, rating, sentiment, ts FROM analysis ORDER BY id DESC LIMIT 5;" Enter password:
+----+-----+-----+-----+-----+
| id | course | instructor | rating | sentiment | ts           |
+----+-----+-----+-----+-----+
| 26078 | Cloud   | Dr. Smith    | 5 | POSITIVE | 2025-08-13 21:24:26 |
| 26073 | Cloud   | John          | 2 | NEUTRAL  | 2025-08-13 19:27:38 |
| 26072 | Ram     | John          | 5 | POSITIVE | 2025-08-13 19:27:10 |
| 26071 | Cloud   | Dr. Smith    | 5 | NEUTRAL  | 2025-08-13 19:06:10 |
| 26070 | Data Science | Prof. Aijith Kumar | 5 | POSITIVE | 2025-08-13 14:25:34 |
+----+-----+-----+-----+-----+
[ec2-user@ip-10-0-137-70 ~]$ mysql -h fmsdb.cgp2uuiquqcg.us-east-1.rds.amazonaws.com -u admin -p -e "USE feedback_raw; SELECT id, course, instructor, rating, LEFT(feedback, 60) AS snippet, ts FROM M_feedbacks ORDER BY id DESC LIMIT 5;" Enter password:
+----+-----+-----+-----+-----+
| id | course | instructor | rating | snippet                | ts           |
+----+-----+-----+-----+-----+
| 26087 | Cloud   | Dr. Smith    | 5 | This is a good test feedback | 2025-08-13 21:24:21 |
| 26086 | Cloud   | Ramesh        | 2 | Bad prof                 | 2025-08-13 19:27:37 |
| 26085 | Ram     | John          | 5 | Great                    | 2025-08-13 19:27:10 |
| 26084 | Cloud   | Dr. Smith    | 5 | This is a test feedback  | 2025-08-13 19:06:10 |
| 26083 | Data Science | Prof. Aijith Kumar | 5 | This is a great program and Prof. Aijith Kumar is a great prof | 2025-08-13 14:25:34 |
+----+-----+-----+-----+-----+
[ec2-user@ip-10-0-137-70 ~]$ exit
logout
Connection to 10.0.137.70 closed.
[ec2-user@ip-10-0-4-230 ~]$ for i in {1..10000}; do curl -s http://54.90.241.60/api/submit-feedback -X POST \
-H "Content-Type: application/json" \
-d '{"student": "Karthi", "course": "Cloud", "instructor": "Prof Ali", "rating": 5, "feedback": "Great"}' >/dev/null &
done

```

The screenshot shows the AWS CloudWatch Metrics Insights interface. A search bar at the top contains the query: `CloudWatch Metrics Insights metrics`. The results table has one row, indicating the metric was last updated 1 minute ago. The table includes columns for Name, Instance ID, Instance state, Status check, Alarm status, Availability Zone, Public IPv4 DNS, and Public IPv4 IP. One instance is listed as Running with a status of 3/3 checks passed. Below the table, a detailed view for the instance `i-0e80f3a597fe2f18b` is shown, specifically for the `WebServer` role. This view includes tabs for Details, Status and alarms, Monitoring, Security, Networking, Storage, and Tags. Under the Details tab, the `Instance summary` section is expanded, showing details like Public IPv4 address (`52.90.161.98`), Instance state (`Running`), Hostname type (`IP name: ip-10-0-4-230.ec2.internal`), and Instance type (`t3.small`). Other tabs show Private IPv4 addresses (`10.0.4.230`), Public DNS (`ec2-52-90-161-98.compute-1.amazonaws.com`), and Elastic IP addresses.

Screenshot of the AWS CloudWatch Activity history for an Auto Scaling group named 'app-asg'.

The activity history shows 40 entries, all of which are successful launches of new EC2 instances. The first few entries are as follows:

- Launching a new EC2 instance: i-02952614ab554eaf5 at 2025-08-13T22:54:54Z due to a monitor alarm TargetTracking-app-asg-AlarmHigh-a12a665b-fff3-45ad-a75a-ad1f0b0f69 in state ALARM triggered policy Target Tracking Policy changing the desired capacity from 2 to 4. At 2025-08-13T22:55:02Z an instance was started in response to a difference between desired and actual capacity, increasing the capacity from 2 to 4.
- Launching a new EC2 instance: i-0e48ce564f2e5036f at 2025-08-13T22:54:54Z due to a monitor alarm Target Tracking app-asg-AlarmHigh-a12a665b-fff3-45ad-a75a-ad1f0b0f69 in state ALARM triggered policy Target Tracking Policy changing the desired capacity from 2 to 4. At 2025-08-13T22:55:02Z an instance was started in response to a difference between desired and actual capacity, increasing the capacity from 2 to 4.
- Launching a new EC2 instance: i-0ebed2302fb8c9c14 at 2025-08-13T22:54:54Z due to a monitor alarm Target Tracking app-asg-AlarmHigh-a12a665b-fff3-45ad-a75a-ad1f0b0f69 in state ALARM triggered policy Target Tracking Policy changing the desired capacity from 1 to 2. At 2025-08-13T22:53:02Z an instance was started in response to a difference between desired and actual capacity, increasing the capacity from 1 to 2.
- Terminating EC2 instance: i-00ud27e3136fe0460 at 2025-08-13T22:54:54Z due to a user request update of AutoScalingGroup constraints to min: 1, max: 1, desired: 1 changing the desired capacity from 4 to 1. At 2025-08-13T14:47:22 an instance was taken out of service in response to a difference between desired and actual capacity, shrinking the capacity from 4 to 1. At 2025-08-13T14:47:22 instance i-0d88ceaa349796688 was selected for termination. At 2025-08-13T14:47:22 instance i-00ud27e3136fe0460 was selected for termination.

The last entry is a terminating instance:

- Terminating EC2 instance: i-00ud27e3136fe0460 at 2025-08-13T22:54:54Z due to a user request update of AutoScalingGroup constraints to min: 1, max: 1, desired: 1 changing the desired capacity from 4 to 1. At 2025-08-13T14:47:22 an instance was taken out of service in response to a difference between desired and actual capacity, shrinking the capacity from 4 to 1. At 2025-08-13T14:47:22 instance i-00ud27e3136fe0460 was selected for termination.

Screenshot of the AWS CloudWatch Targets for a target group named 'app-tg-5000'.

The target group configuration includes:

- Protocol & Port:** HTTP: 5000
- Protocol Version:** HTTP1
- IP address type:** IPv4
- Load balancer:** Internal-App-ALB
- Targets:** 4 total targets, 4 healthy, 0 unhealthy, 0 unused, 0 initial, 0 draining.

The registered targets table lists four targets:

Instance ID	Name	Port	Zone	Health status	Health status details	Administrative o...	Override details	Launch...	Anomaly...
i-02952614ab554eaf5	AppServer-ASG	5000	us-east-1a (us...)	Healthy	-	No override	No override is curre...	August 13...	Normal
i-0e48ce564f2e5036f	AppServer-ASG	5000	us-east-1b (us...)	Healthy	-	No override	No override is curre...	August 13...	Normal
i-0ebed2302fb8c9c14	AppServer-ASG	5000	us-east-1b (us...)	Healthy	-	No override	No override is curre...	August 13...	Normal
i-0d0693f00165b020	AppServer-ASG	5000	us-east-1a (us...)	Healthy	-	No override	No override is curre...	August 13...	Normal

Screenshot of the AWS CloudWatch Metrics console showing the distribution of targets by Availability Zone (AZ) for a Target group named "WebServerTG".

Total targets: 1

Status	Count
Healthy	1
Unhealthy	0
Unused	0
Initial	0
Draining	0

Distribution of targets by Availability Zone (AZ):

Zone	Count
us-east-1a (us...)	1

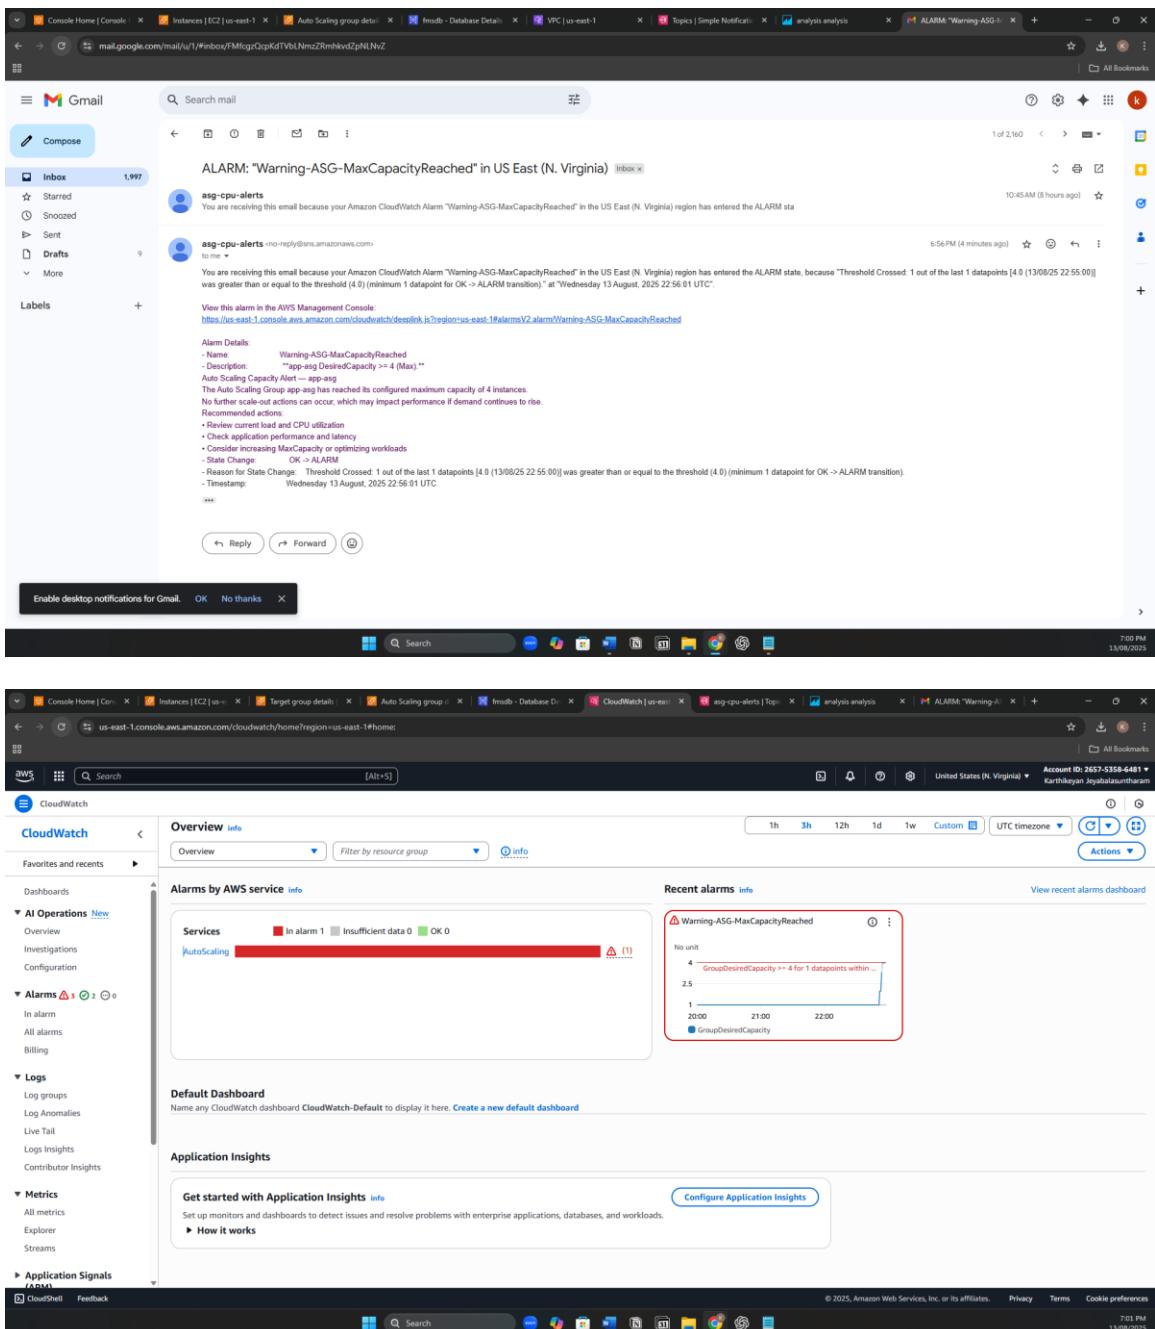
Registered targets (1)

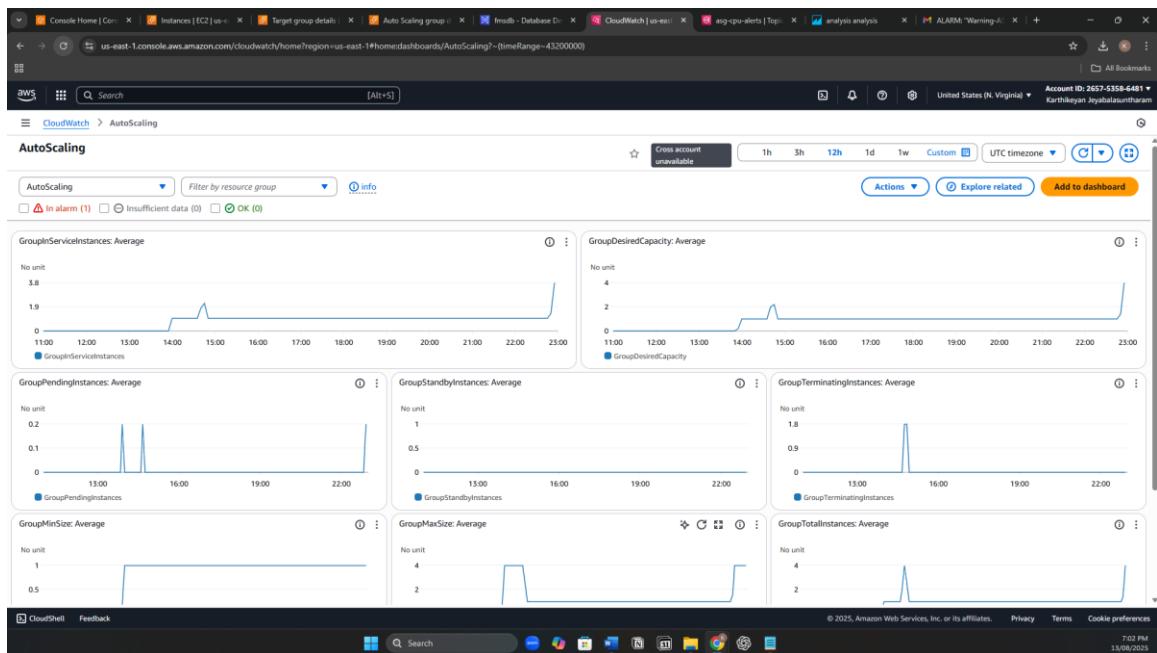
Instance ID	Name	Port	Zone	Health status	Health status details
i-0e0f5a597fe2f18b	WebServer	80	us-east-1a (us...)	Healthy	-

Screenshot of the AWS CloudWatch Metrics console showing the Activity history for an Auto Scaling group named "web-asg".

Activity history (3):

Status	Description	Cause	Start time
Successful	Launching a new EC2 instance: i-00858754ae469927a	At 2025-08-14T20:05:58Z a user request update of AutoScalingGroup constraints to min: 1, max: 4, desired: 1 changing the desired capacity from 0 to 1. At 2025-08-14T20:06:07Z an instance was started in response to a difference between desired and actual capacity, increasing the capacity from 0 to 1.	2025 August 14, 04:06:09 PM -04:00
Successful	Terminating EC2 instance: i-0130b4a6345518056	At 2025-08-14T03:22:38Z a user request update of AutoScalingGroup constraints to min: 0, max: 0, desired: 0 changing the desired capacity from 1 to 0. At 2025-08-14T03:22:47Z an instance was taken out of service in response to a difference between desired and actual capacity, shrinking the capacity from 1 to 0. At 2025-08-14T03:22:47Z instance i-0130b4a6345518056 was selected for termination.	2025 August 13, 11:22:47 PM -04:00
Successful	Launching a new EC2 instance: i-0130b4a6345518056	At 2025-08-13T19:19:57Z a user request created an AutoScalingGroup changing the desired capacity from 0 to 1. At 2025-08-13T19:20:01Z an instance was started in response to a difference between desired and actual capacity, increasing the capacity from 0 to 1.	2025 August 13, 03:20:03 PM -04:00





Web Page

This screenshot shows a 'Course Feedback Form' on a web page. The form consists of the following fields:

- Your Name (Optional)**: Input field containing "Vijay".
- Course Name**: Input field containing "Cloud Computing".
- Instructor Name**: Input field containing "Prof Ajith".
- Rating**: A 5-star rating icon.
- Your Feedback**: Text area containing "This is a great course."
- Submit Feedback**: A blue button at the bottom of the form.

The browser interface at the top shows tabs for 'Course Feedback Form' and 'karthicode.com'. The taskbar at the bottom shows various application icons.

Course Feedback Form

Your Name (Optional)
John Doe

Course Name
Cloud Computing 101

Instructor Name
Dr. Smith

Rating
★ ★ ★ ★ ★

Your Feedback
Write your feedback here...

Submit Feedback

Thank you! Your feedback has been submitted. Sentiment:
POSITIVE

4:13 PM
14/06/2023

Course Feedback Form

Your Name (Optional)
Ramesh

Course Name
Data Science

Instructor Name
Dr. Mark

Rating
★ ★ ★ ★ ★

Your Feedback
This is not a good course.

Submit Feedback

Thank you! Your feedback has been submitted. Sentiment:
POSITIVE

4:15 PM
14/06/2023

Screenshot of a Course Feedback Form application running in a browser window:

Course Feedback Form

Your Name (Optional)
John Doe

Course Name
Cloud Computing 101

Instructor Name
Dr. Smith

Rating
★ ★ ★ ★

Your Feedback
Write your feedback here...

Submit Feedback

Thank you! Your feedback has been submitted. Sentiment:
NEGATIVE

Below the browser window is a terminal session showing MySQL queries and their results:

```

[ec2-user@ip-10-0-137-70:~]
+-----+
| ec2-user@ip-10-0-137-70 ~] $ mysql -h fmsdb.cgpuuiqugcg.us-east-1.rds.amazonaws.com -u admin -p -e "USE feedback_raw; SELECT id, course, instructor, rating, LEFT(feedback,60) AS snippet, ts FROM feedBacks ORDER BY id DESC LIMIT 5;"*
Enter password:
+-----+
| id | course | instructor | rating | snippet | ts
+-----+
| 26092 | Data Science | Mark | 5 | Great | 2025-08-14 20:09:59 |
| 26091 | Cloud Computing | Dr. Smith | 5 | This is a good test feedback | 2025-08-14 20:07:59 |
| 26090 | Data Science | Prof. Mark | 5 | This is a great program. | 2025-08-14 03:08:37 |
| 26089 | Cloud Computing | Dr. Smith | 5 | This is a good test feedback | 2025-08-14 03:05:53 |
| 26088 | Cloud Computing | Dr. Smith | 5 | This is a good test feedback | 2025-08-13 22:53:21 |
+-----+
[ec2-user@ip-10-0-137-70 ~] $ mysql -h fmsdb.cgpuuiqugcg.us-east-1.rds.amazonaws.com -u admin -p -e "USE feedback_raw; SELECT id, course, instructor, rating, LEFT(feedback,60) AS snippet, ts FROM feedBacks ORDER BY id DESC LIMIT 5;"*
Enter password:
+-----+
| id | course | instructor | rating | snippet | ts
+-----+
| 26095 | Data Science | Dr. Mark | 2 | This is not a good course. | 2025-08-14 20:15:26 |
| 26094 | Cloud Computing | Prof. Ajith | 5 | This is a great course. | 2025-08-14 20:13:45 |
| 26093 | Cloud Computing | Prof. Vijay | 5 | This is a great course. | 2025-08-14 20:12:17 |
| 26092 | Data Science | Mark | 5 | Great | 2025-08-14 20:09:59 |
| 26091 | Cloud Computing | Dr. Smith | 5 | This is a good test feedback | 2025-08-14 20:07:59 |
+-----+
[ec2-user@ip-10-0-137-70 ~] $ mysql -h feedbackanalysisdb.cgpuuiqugcg.us-east-1.rds.amazonaws.com -u admin -p -e "USE feedback_analysis; SELECT id, course, instructor, rating, sentiment, ts FROM analysis ORDER BY id DESC LIMIT 5;"*
Enter password:
+-----+
| id | course | instructor | rating | sentiment | ts
+-----+
| 26082 | Data Science | Dr. Mark | 2 | NEGATIVE | 2025-08-14 20:15:26 |
| 26081 | Cloud Computing | Prof. Ajith | 5 | POSITIVE | 2025-08-14 20:13:45 |
| 26080 | Cloud Computing | Prof. Vijay | 5 | POSITIVE | 2025-08-14 20:12:17 |
| 26079 | Data Science | Mark | 5 | POSITIVE | 2025-08-14 20:09:59 |
| 26078 | Cloud Computing | Dr. Smith | 5 | POSITIVE | 2025-08-14 20:07:59 |
+-----+
[ec2-user@ip-10-0-137-70 ~] $
```

Conclusion

The AWS Feedback Management System demonstrates how a modern, cloud-native application can combine scalability, security, and intelligence to deliver real business value. By leveraging a three-tier architecture with Auto Scaling Groups, the system adapts seamlessly to changes in traffic, always ensuring consistent performance while optimizing costs.

Integration with AWS Comprehend enables real-time sentiment analysis, transforming raw user feedback into actionable insights. Amazon QuickSight provides interactive dashboards that make it easy to monitor trends and identify areas for improvement. The use of Amazon SNS for Auto Scaling notifications ensures that operational teams are immediately informed of scaling events, while Amazon CloudWatch maintains continuous visibility into the system's health.